

Multi Agent Systems

- Lab 7 -

Q-Learning with Linear Value Function Approximation

Q-Learning Recap

- Value Function is more explicit in storing the value of executing **an action** in a **given state**: $q(s, a)$

$$q^{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s, A_t = a] = E_{\pi}\left[\sum_{\tau=t+1} \gamma^{\tau-t-1} R_{\tau} | S_t = s, A_t = a\right]$$

- Instance of *model-free learning* – i.e. environment dynamics is unknown to the agent
- We tackled environments where number of states is small enough to use a **tabular** representation of the Q-Function

Q-Learning Recap

Agent learns by observing consequences of actions it takes in the environment

Q-values adjusted through **temporal differences**

Learning is **off-policy**

Learning policy is greedy

Play policy allows for exploration

```
procedure  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
  with prob  $\epsilon$ : return  $\text{random}(A)$   
  with prob  $1-\epsilon$ : return  $\underset{a}{\operatorname{argmax}} q(s, a)$   
end
```

```
procedure Q-Learning ( $\langle S, A, \gamma \rangle, \epsilon$ )  
  for all  $s$  in  $S$ ,  $a$  in  $A$  do  
     $q(s, a) \leftarrow 0$  // set initial values to 0  
  end for  
  for all episodes do  
     $s \leftarrow$  initial state  
    while  $s$  not final state do  
      pick action  $a$  using  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
      execute  $a \rightarrow$  get reward  $r$  and next state  $s'$   
       $q(s, a) \leftarrow q(s, a) + \alpha(r + \gamma \max_{a'} q(s', a') - q(s, a))$   
       $s \leftarrow s'$   
    end while  
  end for  
  for all  $s$  in  $S$  do  
     $\pi(s) \leftarrow \operatorname{argmax}_{a \in A} q(s, a)$   
  end for  
  return  $\pi$ 
```

Q-Learning in continuous state space

- Many real world problems have enormous state and/or action spaces (e.g. robotics control, self driving)
- Tabular representation is not really appropriate
- Idea: Use a function to represent the value

Q-Learning with Linear Value Function Approximation – General Formulation

- Use *features* to represent state and action $x(s, a) = \begin{pmatrix} x_1(s, a) \\ x_2(s, a) \\ \dots \\ x_n(s, a) \end{pmatrix}$

- Q-function represented as ***weighted linear combination of features***

$$\hat{Q}(s, a, \mathbf{w}) = \mathbf{x}(s, a)^T \mathbf{w} = \sum_{j=1}^n x_j(s, a) w_j$$

- ***Learn*** weights $\mathbf{w}$ through stochastic gradient descent updates

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \nabla_{\mathbf{w}} E_{\pi}[(Q^{\pi}(s, a) - \hat{Q}^{\pi}(s, a, \mathbf{w}))^2]$$

Q-Learning with Linear Value Function Approximation – Simplified

- When action space **A** is *small* and *finite*
consider a featurised representation of states only

$$\mathbf{x}(s) = \begin{pmatrix} x_1(s) \\ x_2(s) \\ \dots \\ x_n(s) \end{pmatrix}$$

- Q-function represented as **collection** of *weighted linear combination of features* – **one model per action**

$$\hat{Q}_a(s, \mathbf{w}) = \mathbf{x}(s)^T \mathbf{w} = \sum_{j=1}^n x_j(s) w_j, \forall a \in A$$

- Learn** weights **w** through stochastic gradient descent updates

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \nabla_{\mathbf{w}} E_{\pi}[(Q^{\pi}(s, a) - \hat{Q}_a^{\pi}(s, \mathbf{w}))^2]$$

Q-Learning with Linear Value Function Approximation – TD Target

- For Q-Function, instead of the actual gain per episode under current policy $Q^\pi(s, a)$ use **TD-target** $r + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w})$
- **Learn** weights $\mathbf{w}$ through stochastic gradient descent updates

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \nabla_{\mathbf{w}} (r + \gamma \max_{a'} \hat{Q}_{a'}(s', \mathbf{w}) - \hat{Q}_a(s, \mathbf{w}))^2$$

$$\Delta \mathbf{w} = \eta (r + \gamma \max_{a'} \hat{Q}_{a'}(s', \mathbf{w}) - \hat{Q}_a(s, \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}_a(s, \mathbf{w}), \quad \eta - \text{SGD learning rate}$$

$$\Delta \mathbf{w} = \eta (r + \gamma \max_{a'} \hat{Q}_{a'}(s', \mathbf{w}) - \hat{Q}_a(s, \mathbf{w})) \mathbf{x}(s)$$

Q-Learning, Linear Approximation, TD target

procedure Q-Learning ($\langle S, A, \gamma \rangle, \epsilon$, estimator)

for all episodes **do**

$s \leftarrow$ initial state

while s not final state **do**

 pick action a using ϵ -Greedy (o_s , estimator, ϵ)

 execute $a \rightarrow$ get reward r and next state s'

$x(s') = \text{featurize}(o_{s'})$

$[\hat{q}_{a1}(s'), \dots, \hat{q}_{am}(s')] = \text{estimator.predict}(x(s'))$

$td_{\text{target}} = r + \gamma \max_a \hat{q}_a(s')$

$\text{estimator.update}(s, a, td_{\text{target}})$

$s \leftarrow s'$

end while

end for

for all s in S **do**

$\pi(s) \leftarrow \text{argmax}_{a \in A} q(s, a)$

end for

return π

Agent learns by observing consequences
of actions it takes in the environment

Q-values adjusted through **temporal
differences**

Learning is **off-policy**

 Learning policy is greedy

 Play policy allows for exploration

procedure ϵ -Greedy (o_s , estimator, ϵ)

$x(s) = \text{featurize}(o_s)$

$\hat{q} = \text{estimator}(x(s))$


 with prob ϵ : return $\text{random}(A)$

 with prob $1-\epsilon$: return $\text{argmax}_a \hat{q}_a(s)$

end

OpenAI Gym Cartpole Environment

- **Cartpole-v1** environment in **Gymnasium**:
 - Objective: keep a pendulum upright for as long as possible
 - 2 actions: left (force = -1), right (force = +1)
 - Reward: +1 for every time step that the pole remains upright
 - Game ends when pole more the 15° from vertical OR cart moves > 2.4 units from center



OpenAI Gym Cartpole Environment

- **Cartpole-v1** environment in **Gymnasium**:
 - Max num steps per episode = 100
 - Use one q-function estimator per action: $\hat{q}_{left}, \hat{q}_{right}$
 - Use a MLP-like feature extractor
 - Sample 4 sets of weights and 4 sets of biases
 - $w_{ij}^k \sim \sqrt{2} \gamma_k N(0, 1), k=1..4, i=1..100, j=1..4 \quad \gamma_k \in \{5.0, 2.0, 1.0, 0.5\}$
 - $b_i^k \sim \text{uniform}(0, 2\pi), k=1..4, i=1..100$
 - **Explore** different activation functions: $\cos(x), \text{sigmoid}(x), \tanh(x)$
 - **Explore** different values of the SGD learning rate (5e-4 is default)
 - **Explore** different values of ϵ – $\epsilon=0.0, \epsilon=0.1, \epsilon=\text{decay}(\text{init}=0.1, \text{min} = 0.001, \text{factor}=0.99)$
 - **Plot** agent learning curves for each case